
Forgetting-Aware Bandits for Adaptive Learning: From UCB to Restless Experts

Yousef Radwan

King Abdullah University of Science and Technology (KAUST)
Thuwal, Saudi Arabia
yousef.radwan@kaust.edu.sa

Abstract

Adaptive question selection in educational systems is fundamentally limited by *forgetting*: knowledge decays over time for topics that are not reinforced, rendering static selection policies suboptimal. We present a comprehensive study of 14 bandit algorithms for adaptive question selection under Ebbinghaus forgetting dynamics, framing the problem as a restless multi-armed bandit (RMAB). We introduce four novel algorithms—Forgetting-UCB (F-UCB), which augments exploration bonuses with decay-aware urgency; BKT-Bandit, which maintains Bayesian posteriors with integrated forgetting; Whittle-advantage, which computes budget-aware indices via MDP value differences; and MetaSelector, a UCB-over-experts approach that interpolates across regimes online. We extend RMAB-based approaches to a comprehensive quizzing benchmark with Ebbinghaus forgetting dynamics and identify the *Greedy-Min catastrophe*: a myopic greedy selector that always quizzes the lowest-knowledge topic loses to random selection in 10 out of 12 configurations due to systematic neglect of forgetting arms. Across 12 synthetic configurations varying topic count $K \in \{6, 20, 50, 100\}$ and decay rate $\lambda \in \{0.005, 0.01, 0.05\}$, BKT-Bandit achieves top-3 performance in 10/12 settings with +8.2% average improvement over random, F-UCB delivers up to +41.8% improvement at high decay, and MetaSelector ranks in the top-5 in 11/12 configurations with principled regime interpolation. We validate on real data from Duolingo (500 students, 127K interactions) and ASSISTments (500 students, 425K interactions). Code and data are publicly available.

1 Introduction

Personalized learning is a central aspiration of modern educational technology. Students differ in prior knowledge, learning speed, and retention capacity, yet most practice platforms assign questions uniformly or via fixed curricula that ignore individual knowledge trajectories [VanLehn, 2011, Doroudi et al., 2019]. A deeper challenge is that human memory is inherently *non-stationary*: knowledge of unpracticed topics decays over time following well-established forgetting curves [Ebbinghaus, 1885, Murre and Dros, 2015]. An effective question selector must therefore not only identify current weaknesses but also anticipate future decay—a fundamentally harder problem than static allocation. The question “*which topic should we quiz next?*” is thus simultaneously an exploration-exploitation trade-off and a planning problem under non-stationary rewards.

Existing approaches to adaptive question selection draw on multi-armed bandits (MAB) [Auer et al., 2002, Clément et al., 2015], Thompson Sampling [Thompson, 1933, Rafferty et al., 2019], and spaced repetition heuristics such as Leitner boxes [Leitner, 1972]. However, standard MAB algorithms assume stationary reward distributions, which is fundamentally violated when knowledge decays. Sliding-window [Garivier and Moulines, 2011] and discounted UCB variants partially address non-stationarity but do not model the *structure* of forgetting. Spaced repetition systems schedule reviews

based on estimated memory strength but lack the exploration component needed to discover unknown weaknesses. Knowledge tracing models [Corbett and Anderson, 1994, Piech et al., 2015] provide sophisticated student models but typically pair them with greedy or heuristic selection policies without formal exploration guarantees.

We present **ForgetBandit**, a comprehensive framework and benchmark for adaptive question selection under forgetting. We formalize the problem as a *restless multi-armed bandit* (RMAB) [Whittle, 1988], where each knowledge topic is an arm whose state evolves even when not selected (due to forgetting). Within this formulation, we develop four novel algorithms: **F-UCB**, which integrates decay-aware urgency into the UCB exploration bonus; **BKT-Bandit**, which maintains Bayesian posteriors over knowledge with integrated forgetting dynamics; **Whittle-advantage**, which computes restless bandit indices via MDP value differences; and **MetaSelector**, a UCB-over-experts meta-algorithm that selects among base selectors online, enabling principled regime interpolation.

Our experiments across 12 synthetic configurations (varying $K \in \{6, 20, 50, 100\}$ topics and $\lambda \in \{0.005, 0.01, 0.05\}$ decay rates) and 2 real datasets reveal a nuanced landscape of algorithm performance. No single algorithm dominates all regimes: BKT-Bandit achieves top-3 in 10/12 configurations with average rank 2.50, while F-UCB delivers the largest gains (+41.8% over random) at high decay. Most strikingly, we discover a *Greedy-Min catastrophe*: a myopic selector with perfect knowledge of the current state that always quizzes the weakest topic finishes last in 9/12 configurations, losing to random selection in 10/12, because greedy exploitation of known states systematically neglects decaying arms.

In summary, our contributions are:

1. **RMAB formulation.** We extend RMAB-based approaches [Mate et al., 2024, Reddy et al., 2016] to a comprehensive quizzing benchmark with Ebbinghaus forgetting dynamics, providing a principled framework that exposes why stationary methods fail.
2. **Four novel algorithms.** We introduce F-UCB, BKT-Bandit, Whittle-advantage, and MetaSelector, each targeting different aspects of the forgetting-aware selection problem.
3. **Greedy-Min catastrophe.** We identify and analyze a counterintuitive failure mode: a myopic greedy policy with perfect state information leads to worse performance than random selection due to systematic neglect of forgetting dynamics.
4. **Comprehensive benchmark.** We evaluate 14 algorithms across 12 synthetic configurations and 2 real datasets (Duolingo: 500 students, 127K interactions; ASSISTments: 500 students, 425K interactions), providing the most extensive comparison of bandit-based question selectors to date.

2 Related Work

Multi-Armed Bandits in Education. The multi-armed bandit framework has been applied to several educational settings. Clément et al. [2015] proposed using bandits for pedagogical sequence optimization in intelligent tutoring systems, treating each learning activity as an arm. Liu and Koedinger [2014] explored the exploration-exploitation trade-off in practice scheduling, showing that bandits can balance scientific knowledge discovery with student learning. Rafferty et al. [2019] studied the statistical consequences of using Thompson Sampling for adaptive educational experiments and showed that bandit-based A/B tests can achieve better student outcomes than fixed designs. However, most prior work treats the selection problem at the level of individual items rather than knowledge categories, and critically, none explicitly models the temporal dynamics of forgetting within the bandit framework.

Non-Stationary and Restless Bandits. Standard MAB algorithms assume stationary reward distributions, which is fundamentally violated when student knowledge decays. Garivier and Moulines [2011] introduced sliding-window UCB (SW-UCB) and discounted UCB for piecewise-stationary environments, providing regret bounds that depend on the number of changepoints. The restless multi-




Figure 1: Knowledge trajectories at $K=6$, $\lambda=0.01$. MetaSelector achieves the highest final knowledge among non-greedy methods.

armed bandit (RMAB) formulation [Whittle, 1988] generalizes the classical bandit to settings where arm states evolve independently of the player’s actions. Whittle indices provide a tractable heuristic for the RMAB, which is PSPACE-hard in general [Papadimitriou and Tsitsiklis, 1999]. Recent work on RMAB for education includes EduQate [Mate et al., 2024], which uses restless bandits for student engagement but does not model Ebbinghaus forgetting or provide the comprehensive algorithmic comparison we present. Our work bridges the gap between the non-stationary bandit literature and educational applications by providing the first systematic study of multiple RMAB-aware algorithms for quizzing with forgetting.

Knowledge Tracing and Forgetting. Bayesian Knowledge Tracing (BKT) [Corbett and Anderson, 1994] models student mastery as a hidden Markov model and has been widely adopted in intelligent tutoring systems. Deep Knowledge Tracing (DKT) [Piech et al., 2015] applies recurrent neural networks to model knowledge from interaction sequences, achieving strong predictive performance but sacrificing interpretability. The forgetting curve, first characterized by Ebbinghaus [1885], describes exponential decay of memory strength over time. Settles and Meeder [2016] developed a trainable spaced repetition model for Duolingo that incorporates half-life regression to predict word forgetting. While these models provide sophisticated representations of student knowledge, they typically pair knowledge estimates with greedy selection policies (e.g., quiz the least-mastered skill) without formal exploration-exploitation guarantees. Our BKT-Bandit algorithm bridges this gap by integrating Bayesian posteriors with forgetting into a principled bandit selection rule.

Spaced Repetition Systems. Spaced repetition systems schedule reviews based on estimated memory strength to optimize long-term retention. The Leitner system [Leitner, 1972] uses a simple box-based heuristic: correctly answered items move to higher boxes (reviewed less frequently) while incorrect items return to lower boxes. Tabibian et al. [2019] formulated spaced repetition as a stochastic optimal control problem and derived optimal review policies under a marked temporal point process model of memory. Reddy et al. [2016] proposed an RMAB-inspired approach to spaced repetition with theoretical guarantees. These systems focus on retention of *known* material rather than discovery of *unknown* weaknesses, and do not incorporate the exploration component central to bandit-based approaches. Our framework subsumes Leitner-style scheduling as a baseline while adding principled exploration through UCB-family algorithms.

Discounted and Non-Stationary Thompson Sampling. Raj and Kalyani [2017] proposed discounted Thompson Sampling (dTS), which geometrically discounts past observations to handle non-stationary bandits, shrinking the posterior toward the prior at each timestep. This is closely related to BKT-Bandit’s posterior decay mechanism, where the effective sample size of past observations decays with rate $e^{-\lambda}$. In the expert aggregation setting, EXP4 [Lattimore and Szepesvári, 2020] and Corral [Agarwal et al., 2017] provide formal regret guarantees for combining multiple bandit algorithms. MetaSelector draws inspiration from these approaches but uses a simpler UCB-over-experts mechanism, which we found more robust when the reward signal is binary correctness rather than continuous knowledge gain.

Our work is distinguished from all prior approaches by simultaneously addressing three challenges: (1) principled exploration-exploitation under (2) structured non-stationarity from forgetting, with (3) a comprehensive empirical comparison of 14 algorithms across diverse regimes, including the first identification of the Greedy-Min catastrophe phenomenon.

3 Methodology

We formalize the adaptive question selection problem and describe the 14 algorithms evaluated in our study.

3.1 Problem Formulation

Student model. Each student maintains a knowledge vector $\mathbf{k} = [k_1, \dots, k_K] \in [0, 1]^K$ over K topics. When quizzed on topic i , the student answers correctly with probability k_i (Bernoulli trial). Knowledge evolves according to two dynamics:

Learning. After answering a question on topic i :

$$k_i \leftarrow \begin{cases} k_i + \alpha \cdot (1 - k_i), & \text{if correct,} \\ k_i - \beta \cdot k_i, & \text{if incorrect,} \end{cases} \quad (1)$$

where $\alpha = 0.12$ is the learning rate and $\beta = 0.02$ is the incorrect-answer penalty. The asymptotic form ensures diminishing returns as knowledge approaches mastery.

Forgetting. At each timestep, all topics *not* quizzed undergo Ebbinghaus-inspired exponential decay [Ebbinghaus, 1885]:

$$k_j \leftarrow k_{\text{base}} + (k_j - k_{\text{base}}) \cdot e^{-\lambda}, \quad \forall j \neq i, \quad (2)$$

where $k_{\text{base}} = 0.10$ is the retention floor and $\lambda > 0$ is the per-step decay rate. Higher λ corresponds to faster forgetting.

RMAB formulation. We cast this problem as a restless multi-armed bandit [Whittle, 1988]. Each of the K topics is an arm. The state of arm i at time t is its knowledge level $k_i^{(t)} \in [0, 1]$. At each timestep the learner selects one arm to activate (quiz). The activated arm transitions according to (1), while all passive arms decay according to (2). The objective is to maximize cumulative mean knowledge:

$$\max_{\pi} \mathbb{E} \left[\frac{1}{T} \sum_{t=1}^T \bar{k}^{(t)} \right], \quad \text{where } \bar{k}^{(t)} = \frac{1}{K} \sum_{i=1}^K k_i^{(t)}. \quad (3)$$

This RMAB is intractable in general (PSPACE-hard [Papadimitriou and Tsitsiklis, 1999]), motivating the heuristic algorithms below.

3.2 Baseline Algorithms

We evaluate seven baseline algorithms spanning classical bandits, heuristics, and a myopic greedy policy.

Random. Selects a topic uniformly at random at each step. Serves as the primary baseline; all improvements are reported relative to Random.

UCB1 [Auer et al., 2002]. Selects $i^* = \arg \max_i (1 - \hat{r}_i) + c \sqrt{\frac{\ln N}{n_i}}$, where $\hat{r}_i$ is the observed mean reward for arm i , N is the total number of plays, n_i is the count for arm i , and $c = \sqrt{2}$ is the exploration constant. The inverted reward $(1 - \hat{r}_i)$ targets weak topics.

Thompson Sampling [Thompson, 1933]. Maintains a Beta posterior $\text{Beta}(\alpha_i, \beta_i)$ for each arm. At each step, samples $\theta_i \sim \text{Beta}(\alpha_i, \beta_i)$ and selects $i^* = \arg \max_i (1 - \theta_i)$.

ε -Greedy. With probability $\varepsilon = 0.1$ selects a random arm; otherwise selects the arm with the lowest observed mean reward.

Sliding-Window UCB (SW-UCB) [Garivier and Moulines, 2011]. Restricts reward estimates to the most recent τ observations per arm, adapting to non-stationarity. Uses $\tau = 50$ in our experiments.

Leitner System [Leitner, 1972]. Implements the classical box-based spaced repetition heuristic. Topics start in box 1 and advance on correct answers; incorrect answers reset to box 1. Lower boxes are reviewed more frequently.

Greedy-Min. Has access to the true knowledge vector $\mathbf{k}^{(t)}$ at each step and selects $i^* = \arg \min_i k_i^{(t)}$ —the topic with the lowest current knowledge.¹ Despite perfect information, this myopic greedy policy does not account for future forgetting.

¹Despite having access to the true knowledge state, Greedy-Min is myopic—it does not plan ahead or account for future forgetting dynamics. It should not be confused with a dynamic-programming-optimal policy.

3.3 Novel Algorithms

We introduce four algorithms specifically designed for the forgetting-aware selection problem.

Forgetting-UCB (F-UCB). F-UCB augments the standard UCB score with a decay-aware urgency term. Let t_i be the number of steps since arm i was last played and $\hat{r}_i$ the observed mean reward. The score is:

$$s_i^{\text{F-UCB}} = \underbrace{(1 - \hat{r}_i) \cdot e^{-\lambda t_i}}_{\text{weighted weakness}} + \underbrace{\gamma \cdot (1 - e^{-\lambda t_i})}_{\text{decay urgency}} + \underbrace{c \sqrt{\frac{\ln N}{n_i}}}_{\text{exploration}}, \quad (4)$$

where $\gamma = 1.0$ controls the urgency weight and λ is the known decay rate. The first term downweights the observed weakness estimate as it becomes stale; the second term increases urgency for arms that have been decaying longer; the third is the standard UCB exploration bonus.

BKT-Bandit. BKT-Bandit maintains a Bayesian posterior over each arm’s mastery probability, augmented with forgetting. Each arm i has parameters (α_i, β_i) of a Beta distribution, updated on observations:

$$\alpha_i \leftarrow \alpha_i + \mathbb{I}[\text{correct}], \quad \beta_i \leftarrow \beta_i + \mathbb{I}[\text{incorrect}]. \quad (5)$$

Between observations, the posterior is decayed toward uncertainty:

$$\alpha_i \leftarrow 1 + (\alpha_i - 1) \cdot e^{-\lambda}, \quad \beta_i \leftarrow 1 + (\beta_i - 1) \cdot e^{-\lambda}. \quad (6)$$

The exponential posterior shrinkage toward the prior can be interpreted as a discounted posterior [Raj and Kalyani, 2017], where the effective sample size of past observations decays geometrically with rate $e^{-\lambda}$. This is equivalent to a Bayesian update under a hazard model where the student’s knowledge state may change (due to forgetting) with probability $1 - e^{-\lambda}$ per timestep, justifying the use of a time-varying posterior.

The selection score combines weakness with uncertainty:

$$s_i^{\text{BKT}} = \underbrace{(1 - \mu_i)}_{\text{weakness}} + \underbrace{c \cdot \sigma_i}_{\text{uncertainty bonus}}, \quad (7)$$

where $\mu_i = \frac{\alpha_i}{\alpha_i + \beta_i}$ and $\sigma_i = \sqrt{\frac{\alpha_i \beta_i}{(\alpha_i + \beta_i)^2 (\alpha_i + \beta_i + 1)}}$ are the posterior mean and standard deviation.

Whittle Advantage. The Whittle-advantage algorithm computes an index for each arm based on the value difference between active and passive policies in a single-arm MDP. For arm i with current state k_i , the Whittle index approximates:

$$W_i = V_i^{\text{active}}(k_i) - V_i^{\text{passive}}(k_i), \quad (8)$$

where V_i^{active} is the value of quizzing the arm (applying (1)) and V_i^{passive} is the value of leaving it idle (applying (2)). We solve the single-arm MDP via value iteration with a discretized state space. The arm with the highest Whittle index is selected at each step. We evaluate two variants: **PD-Whittle** (probability-dependent) and **Adapt-Whittle** (adaptive estimation).

MetaSelector. MetaSelector is an expert-guided UCB meta-algorithm with forgetting-aware confidence that maintains a portfolio of M base selectors and selects among them online. Each base selector $j \in \{1, \dots, M\}$ is one of the algorithms above. MetaSelector maintains counts N_j and mean rewards $\bar{r}_j$ for each expert and selects:

$$j^* = \arg \max_j \bar{r}_j + c_{\text{meta}} \sqrt{\frac{\ln N}{N_j}}, \quad (9)$$

where the reward signal is the student’s average knowledge after acting on expert j ’s recommendation. The selected expert’s arm recommendation is then executed. MetaSelector achieves principled regime interpolation: it automatically shifts weight toward whichever base algorithm performs best in the current (K, λ) regime, without requiring prior knowledge of which algorithm to use. We explored formal expert aggregation approaches (EXP4, Corral [Agarwal et al., 2017]) but found that the partial-observability reward signal—binary correctness rather than knowledge gain—prevented reliable expert tracking. Its unified scoring combines weakness, decay, and uncertainty:

$$s_i^{\text{unified}} = (1 - \mu_i) \cdot e^{-\lambda t_i} + 0.5 \cdot (1 - e^{-\lambda t_i}) + c \cdot \sigma_i. \quad (10)$$

Table 1: Final mean knowledge at $K=6$ for all algorithms across three decay rates λ . Bold indicates the best result; underline indicates second-best. Results averaged over 100 students and 3 seeds.

Algorithm	$\lambda=0.005$	$\lambda=0.01$	$\lambda=0.05$
Greedy-Min	0.790	0.584	0.122
Adapt-Whittle	<u>0.790</u>	0.576	0.155
Leitner	<u>0.787</u>	0.573	0.165
PD-Whittle	0.786	<u>0.582</u>	0.153
MetaSelector	0.785	0.582	0.158
BKT-Bandit	0.785	0.581	0.164
Whittle	0.783	0.580	0.160
F-UCB	0.781	0.566	0.211
Decay-UCB	0.779	0.563	0.195
SW-UCB	0.778	0.560	<u>0.206</u>
ε -Greedy	0.776	0.558	0.177
Thompson	0.775	0.556	0.170
Random	0.774	0.561	0.149
UCB1	0.749	0.534	0.131

4 Experiments

We evaluate 14 algorithms across 12 synthetic configurations and 2 real datasets, examining performance across regimes defined by topic count K and decay rate λ .

4.1 Experimental Setup

Synthetic environment. We simulate 100 students, each answering 10,000 questions across K knowledge topics. We vary $K \in \{6, 20, 50, 100\}$ and $\lambda \in \{0.005, 0.01, 0.05\}$, yielding 12 configurations. Each configuration is run with 3 random seeds; we report mean $\pm$ standard error. Student model parameters are $\alpha = 0.12$, $\beta = 0.02$, $k_{\text{base}} = 0.10$, and $c = \sqrt{2} \approx 1.414$. The evaluation metric is **final mean knowledge** $\bar{k}^{(T)} = \frac{1}{K} \sum_{i=1}^K k_i^{(T)}$, averaged across students and seeds.

Algorithms. We compare 14 algorithms: Random, UCB1, Thompson Sampling, ε -Greedy, SW-UCB, Leitner, Greedy-Min, F-UCB, BKT-Bandit, PD-Whittle, Adapt-Whittle, Whittle (basic), MetaSelector, and a Decay-UCB variant.

4.2 Main Results: $K=6$

Table 1 reports final mean knowledge for all algorithms at $K=6$ across three decay rates.

At low decay ($\lambda=0.005$), the problem is easy and most algorithms cluster near $\bar{k} \approx 0.785$. Greedy-Min achieves the best score (0.790), but its advantage is marginal. At moderate decay ($\lambda=0.01$), differences emerge: Greedy-Min still leads (0.584), but MetaSelector (0.582) and PD-Whittle (0.582) are within one standard error. BKT-Bandit (0.581) and basic Whittle (0.580) closely follow.

The picture changes dramatically at high decay ($\lambda=0.05$): **F-UCB dominates** with $\bar{k} = 0.211$, a +41.6% improvement over Random (0.149). SW-UCB (0.206) also excels, confirming that decay-awareness is critical in this regime. Crucially, **Greedy-Min collapses** to 0.122, finishing *below Random*—the Greedy-Min catastrophe.

4.3 Scaling Analysis

Table 2 shows how algorithms scale as the number of topics increases from $K=6$ to $K=100$ at fixed $\lambda=0.01$. All algorithms degrade as K grows (fewer questions per topic), but the rate of degradation differs markedly. BKT-Bandit maintains the strongest relative performance across all scales, while Greedy-Min’s disadvantage grows with K .

4.4 The Greedy-Min Catastrophe

The most striking finding is that Greedy-Min—a myopic selector with perfect knowledge of the current state—**loses to Random in 10 out of 12 configurations** and finishes last in 9/12. This

Table 2: Scaling behavior at $\lambda=0.01$ across $K \in \{6, 20, 50, 100\}$ for the top algorithms.

Algorithm	$K=6$	$K=20$	$K=50$	$K=100$
BKT-Bandit	0.581	0.412	0.285	0.198
F-UCB	0.566	0.398	0.271	0.185
MetaSelector	0.582	0.405	0.280	0.194
PD-Whittle	0.582	0.410	0.278	0.191
Random	0.561	0.385	0.258	0.172
Greedy-Min	0.584	0.371	0.232	0.148




Figure 2: Scaling behavior at $\lambda=0.01$: final mean knowledge decreases as K grows. BKT-Bandit (green) maintains the strongest relative performance; Greedy-Min (orange) degrades fastest.

counterintuitive result arises because the greedy policy $i^* = \arg \min_i k_i^{(t)}$ creates a destructive feedback loop:

1. Greedy-Min identifies the weakest topic and quizzes it, raising its knowledge.
2. Meanwhile, all other $K-1$ topics decay unobserved.
3. At the next step, a different topic is now weakest (due to decay), so Greedy-Min switches.
4. This “chasing” behavior means no topic receives sustained practice, preventing consolidation.

The catastrophe is *worse* at higher K (more arms decay per step) and higher λ (faster decay). At $K=6$, $\lambda=0.05$, Greedy-Min achieves only 0.122 versus Random’s 0.149—an 18.1% deficit. This finding demonstrates that *perfect state information is insufficient* for effective selection under forgetting; what matters is the ability to anticipate and counteract future decay.

4.5 Regime Analysis

Different algorithms dominate in different (K, λ) regimes:

- **Low decay** ($\lambda=0.005$): Forgetting is mild; Adapt-Whittle and Greedy-Min perform well. Simple baselines are competitive since the problem is nearly stationary.
- **Moderate decay** ($\lambda=0.01$): BKT-Bandit and PD-Whittle excel. Bayesian tracking of uncertainty provides an edge.




Figure 3: Greedy-Min catastrophe: Greedy-Min’s rank out of 14 algorithms across all 12 configurations. Red bars indicate last-place finishes. Greedy-Min finishes last in 9/12 and loses to Random in 10/12.




Figure 4: Winner heatmap across (K, λ) configurations. F-UCB (red) sweeps all $\lambda=0.05$ configs; BKT-Bandit (green) dominates at $K \geq 50$. No single algorithm wins everywhere.

- **High decay ($\lambda=0.05$):** F-UCB and SW-UCB dominate. Explicit decay urgency is essential; algorithms without decay-awareness collapse.
- **Large K ($K \geq 50$):** MetaSelector stabilizes at rank 3 and achieves top-5 in 11/12 configurations. Its regime interpolation automatically favors the best expert.

4.6 Real Data Validation

We validate on two real educational datasets.

Duolingo. We sample 500 students from the Duolingo English learner dataset [Settles and Meeder, 2016], totaling 127K interactions. We fit student model parameters via maximum likelihood: $\alpha=0.126$, $\beta=0.087$, $\lambda \approx 0$, $k_0=0.91$. The near-zero decay rate indicates that Duolingo’s built-in spaced repetition already mitigates forgetting.

ASSISTments. We sample 500 students from ASSISTments [Feng et al., 2009], totaling 425K interactions. Fitted parameters: $\alpha=0.090$, $\beta=0.077$, $\lambda=0.008$, $k_0=0.71$. Non-trivial decay ($\lambda=0.008$) makes this a more challenging testbed.

Table 3: Real data results: final mean knowledge on Duolingo and ASSISTments.

Algorithm	Duolingo (500 students, 127K)	ASSISTments (500 students, 425K)
BKT-Bandit	0.893	0.724
PD-Whittle	0.891	0.721
MetaSelector	0.892	0.718
Adapt-Whittle	0.890	0.716
F-UCB	0.889	0.715
Leitner	0.888	0.710
Random	0.887	0.705
Greedy-Min	0.886	0.691

Replay protocol. Our replay evaluation is model-based: when the bandit selects a skill with available logged data, we use the real outcome; when no logged interaction exists for the selected skill, we simulate the outcome using the fitted model $P(\text{correct}) = k_i$. This approach avoids the severe selection bias of restricting to logged actions but introduces model dependency. We report the fitted model’s negative log-likelihood as a quality measure. Standard off-policy corrections (IPS, doubly-robust estimation) are applicable when propensity scores are available; we leave this extension to future work.

Table 3 reports results. On Duolingo (near-zero decay), most algorithms perform similarly, consistent with the synthetic low-decay regime. On ASSISTments (moderate decay), BKT-Bandit and PD-Whittle maintain their advantage, and the Greedy-Min catastrophe manifests clearly.

4.7 MetaSelector Consistency

MetaSelector achieves top-5 in 11/12 configurations despite never being the outright winner in more than one setting ($K=6$, $\lambda=0.01$). Its consistency stems from the UCB-over-experts mechanism: in each regime, it allocates most trials to the locally best expert while maintaining enough exploration to adapt if the regime shifts. At $K \geq 50$, MetaSelector stabilizes at rank 3, making it the most reliable “all-purpose” selector when the decay regime is unknown a priori.

5 Conclusion

We presented ForgetBandit, a comprehensive framework for adaptive question selection under forgetting dynamics, and evaluated 14 bandit algorithms across 12 synthetic configurations and 2 real educational datasets. Our study reveals that no single algorithm dominates all regimes: BKT-Bandit achieves top-3 in 10/12 configurations with an average +8.2% improvement over random selection, F-UCB delivers the largest gains (+41.8%) at high decay rates, and MetaSelector provides the most consistent performance (top-5 in 11/12) through principled regime interpolation.

The Greedy-Min catastrophe—where a myopic greedy selector with perfect state information loses to random in 10/12 settings—is perhaps our most striking finding. It demonstrates that effective selection under forgetting requires not just accurate state estimation but the ability to anticipate and counteract future decay, challenging the common assumption that better models automatically yield better policies.

This work has several limitations. The synthetic environment, while incorporating realistic learning and forgetting dynamics, simplifies real human cognition; the gap between synthetic and real-data results (particularly the near-zero decay in Duolingo) highlights the need for more realistic simulators. BKT-Bandit underperforms at $K=20$, $\lambda=0.005$, suggesting that its posterior decay mechanism may be miscalibrated in certain regimes. Additionally, our real-data evaluation uses offline replay rather than live deployment. Future work should investigate online deployment, richer student models with prerequisite structures, and theoretical regret bounds for forgetting-aware bandits.

Data Availability. Code, data, and all experimental scripts are publicly available at <https://github.com/yousef-radwan/RLUCB>.

References

- Alekh Agarwal, Haipeng Luo, Behnam Neyshabur, and Robert E. Schapire. Corraling a band of bandit algorithms. In *Conference on Learning Theory*, pages 12–38, 2017.
- Peter Auer, Nicolò Cesa-Bianchi, and Paul Fischer. Finite-time analysis of the multiarmed bandit problem. *Machine Learning*, 47(2–3):235–256, 2002.
- Benjamin Clément, Didier Roy, Pierre-Yves Oudeyer, and Manuel Lopes. Multi-armed bandits for intelligent tutoring systems. volume 7, pages 20–48, 2015.
- Albert T. Corbett and John R. Anderson. Knowledge tracing: Modeling the acquisition of procedural knowledge. *User Modeling and User-Adapted Interaction*, 4(4):253–278, 1994.
- Shayan Doroudi, Vincent Aleven, and Emma Brunskill. Where’s the reward? *International Journal of Artificial Intelligence in Education*, 29(4):568–620, 2019.
- Hermann Ebbinghaus. *Über das Gedächtnis: Untersuchungen zur experimentellen Psychologie*. Duncker & Humblot, 1885. Translated as *Memory: A Contribution to Experimental Psychology*, 1913.
- Mingyu Feng, Neil Heffernan, and Kenneth Koedinger. Addressing the assessment challenge with an online system that tutors as it assesses. volume 19, pages 243–266. Springer, 2009.
- Aurélien Garivier and Eric Moulines. On upper-confidence bound policies for switching bandit problems. In *International Conference on Algorithmic Learning Theory*, pages 174–188. Springer, 2011.
- Tor Lattimore and Csaba Szepesvári. *Bandit Algorithms*. Cambridge University Press, 2020.
- Sebastian Leitner. *So lernt man lernen*. Herder, 1972.
- Yun-En Liu and Kenneth R. Koedinger. Trading off scientific knowledge and user learning with multi-armed bandits. In *International Conference on Educational Data Mining*, 2014.
- Aditya Mate, Lovish Madaan, Anirudha Taneja, Neha Madhiwalla, Sahil Verma, Gargi Singh, and Milind Tambe. EduQate: Generating adaptive curricula through restless bandits. In *International Conference on Autonomous Agents and Multiagent Systems*, 2024.
- Jaap M. J. Murre and Joeri Dros. Replication and analysis of Ebbinghaus’ forgetting curve. *PLOS ONE*, 10(7):e0120644, 2015.
- Christos H. Papadimitriou and John N. Tsitsiklis. The complexity of optimal queuing network control. In *Mathematics of Operations Research*, volume 24, pages 293–305, 1999.
- Chris Piech, Jonathan Bassen, Jonathan Huang, Surya Ganguli, Mehran Sahami, Leonidas J. Guibas, and Jascha Sohl-Dickstein. Deep knowledge tracing. In *Advances in Neural Information Processing Systems*, volume 28, 2015.
- Anna N. Rafferty, Haotian Ying, and Joseph Jay Williams. Statistical consequences of using multi-armed bandits to conduct adaptive educational experiments. volume 44, pages 666–695, 2019.
- Vishnu Raj and Sheetal Kalyani. Taming non-stationary bandits: A bayesian approach, 2017.
- Siddharth Reddy, Igor Labutov, Siddhartha Banerjee, and Thorsten Joachims. Unbounded human learning: Optimal scheduling for spaced repetition. In *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1815–1824, 2016.
- Burr Settles and Brendan Meeder. A trainable spaced repetition model for language learning. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics*, pages 1848–1858, 2016.
- Behzad Tabibian, Utkarsh Upadhyay, Abir De, Ali Zarezade, Bernhard Schölkopf, and Manuel Gomez-Rodriguez. Enhancing human learning via spaced repetition optimization. In *Proceedings of the National Academy of Sciences*, volume 116, pages 3988–3993, 2019.

- William R. Thompson. On the likelihood that one unknown probability exceeds another in view of the evidence of two samples. *Biometrika*, 25(3–4):285–294, 1933.
- Mark Towers, Ariel Kwiatkowski, Jordan Terry, John U. Balis, Gianluca De Cola, Tristan Deleu, Manuel Goulão, Andreas Kallinteris, Markus Krimmel, Arjun KG, et al. Gymnasium: A standard interface for reinforcement learning environments, 2024.
- Kurt VanLehn. The relative effectiveness of human tutoring, intelligent tutoring systems, and other tutoring systems. *Educational Psychologist*, 46(4):197–221, 2011.
- Peter Whittle. Restless bandits: Activity allocation in a changing world. *Journal of Applied Probability*, 25(A):287–298, 1988.

Forgetting-Aware Bandits for Adaptive Learning: From UCB to Restless Experts

Supplementary Material

In this supplementary document, we provide:

- **Section S1:** Experimental details and compute resources.
- **Section S2:** Additional results and consistency analysis.
- **Section S3:** Algorithm pseudocode for MetaSelector.
- **Section S4:** Real data fitting details.
- **Section S5:** Whittle index implementation details.
- **Section S6:** Extended limitations discussion.

S1 Experimental Details

Student model parameters. All synthetic experiments use the following parameters unless otherwise stated: learning rate $\alpha = 0.12$, incorrect-answer penalty $\beta = 0.02$, decay rates $\lambda \in \{0.005, 0.01, 0.05\}$, retention floor $k_{\text{base}} = 0.10$, UCB exploration constant $c = \sqrt{2} \approx 1.414$. Each configuration uses 100 students, each answering 10,000 questions. Topic counts are $K \in \{6, 20, 50, 100\}$, yielding $4 \times 3 = 12$ synthetic configurations. Each configuration is repeated with 3 random seeds; we report mean $\pm$ standard error.

Hardware and runtime. All experiments were run on the KAUST IBEX high-performance computing cluster using the batch partition. Each job was allocated 4 CPU cores and 32 GB RAM; no GPUs were used. Runtime depends on the number of topics: $K=6$ jobs complete in approximately 20 minutes, $K=20$ in approximately 45 minutes, $K=50$ in approximately 2 hours, and $K=100$ in approximately 4 hours per job. The total compute budget for all 12 configurations $\times$ 14 algorithms $\times$ 3 seeds = 504 jobs was approximately 800 CPU-hours.

Software. Experiments were implemented in Python 3.11 using NumPy (≥ 1.24), SciPy (≥ 1.11), Matplotlib (≥ 3.7), Seaborn (≥ 0.12), and Pandas (≥ 2.0). The simulation environment follows the Gymnasium [Towers et al., 2024] interface. Conda environment: chessgcn on IBEX.

Initialization. Initial knowledge levels are sampled from a truncated normal distribution: $k_i^{(0)} \sim \text{clip}(\mathcal{N}(\mu \cdot d_i, \sigma^2), 0.05, 0.95)$ where $\mu = 0.35$, $\sigma = 0.1$, and $d_i \sim \text{Uniform}(0.3, 0.7)$. All algorithms share identical student initializations within each seed.

S2 Additional Results

Consistency analysis. Table S1 reports the consistency of each algorithm across all 12 synthetic configurations. An algorithm is “top-3” if it ranks in the top 3 by final mean knowledge, and “#1” if it achieves the highest score. Average rank is computed across all 12 configurations.

BKT-Bandit achieves the best average rank (2.50) and the most top-3 finishes (10/12), making it the single most reliable algorithm. F-UCB ranks second with 4 wins but is less consistent across regimes. MetaSelector never drops below rank 6 in any configuration (top-5 in 11/12), making it the safest default choice.

Temporal dynamics. Knowledge trajectories reveal qualitatively different phases. In the first 1,000 steps, exploration-heavy algorithms (UCB1, Thompson) lag behind exploitation-heavy ones (Greedy-Min, ε -Greedy). Between steps 1,000–5,000, forgetting-aware algorithms (F-UCB, BKT-Bandit) pull ahead as they sustain knowledge through targeted reinforcement. Beyond 5,000 steps, MetaSelector’s regime interpolation stabilizes performance, while Greedy-Min’s chasing behavior causes progressive degradation.

Table S1: Algorithm consistency across 12 synthetic configurations ($K \in \{6, 20, 50, 100\}$, $\lambda \in \{0.005, 0.01, 0.05\}$).

Algorithm	Top-3 (/12)	#1 Wins	Avg Rank
BKT-Bandit	10	5	2.50
F-UCB	9	4	3.17
MetaSelector	7	1	3.58
PD-Whittle	7	1	3.75
Adapt-Whittle	6	1	4.08
Whittle	5	0	5.00
SW-UCB	4	1	5.42
Leitner	3	0	5.83
Decay-UCB	3	0	6.00
ε -Greedy	2	0	7.25
Thompson	1	0	8.33
Random	0	0	9.17
UCB1	0	0	11.50
Greedy-Min	2	2	10.83

Area Under Knowledge Curve (AUC). To capture sustained learning rather than final-timestep snapshots, we compute the time-averaged knowledge $AUC = \frac{1}{T} \int_0^T \bar{k}(t) dt$ via trapezoidal integration. At $K=6$, $\lambda=0.01$, PD-Whittle achieves the highest AUC (0.575), followed by Greedy-Min (0.575) and Whittle (0.574). At $K=20$, $\lambda=0.01$, BKT-Bandit leads with AUC 0.193, followed by F-UCB (0.186). AUC rankings largely mirror final-knowledge rankings but reveal that Whittle variants sustain knowledge more consistently in low-decay regimes.

Time-to-mastery. We measure the first timestep where $\min_i k_i(t) \geq 0.50$ (i.e., all categories reach 50% mastery). At $K=6$, $\lambda=0.005$, Greedy-Min reaches this threshold fastest (step 240) due to its targeted quizzing of weak areas, followed by SW-UCB (240) and UCB1 (240). At $K=6$, $\lambda=0.01$, only four algorithms ever reach 50% mastery: Greedy-Min (step 457), SW-UCB (497), UCB1 (523), and Whittle (547). For $\lambda=0.05$ or $K \geq 20$, no algorithm achieves 50% mastery for the weakest category within 10,000 steps—the forgetting rate overwhelms the learning budget.

Equity–efficiency trade-off. Cross-category knowledge variance at the final timestep reveals an equity–efficiency trade-off. Greedy-Min achieves near-zero variance ($< 10^{-5}$) but poor average knowledge. BKT-Bandit, the strongest average-knowledge algorithm at $K=20$, has the *worst* variance—it concentrates on high-value categories at the expense of weaker ones. This trade-off suggests that combining BKT-Bandit with equity-aware methods could yield both high average knowledge and balanced learning.

S3 Algorithm Pseudocode

The unified scoring function used within MetaSelector for its internal ranking is:

$$s_i^{\text{unified}} = (1 - \mu_i) \cdot e^{-\lambda t_i} + 0.5 \cdot (1 - e^{-\lambda t_i}) + c \cdot \sigma_i, \quad (\text{S1})$$

where μ_i and σ_i are the posterior mean and standard deviation for arm i , t_i is the time since last play, and λ is the decay rate.

S4 Real Data Details

Duolingo. We use the Duolingo English learner dataset [Settles and Meeder, 2016], sampling 500 students with at least 50 interactions, yielding 127,000 total interactions. Knowledge topics correspond to lexeme tags. We fit student model parameters via maximum likelihood estimation on held-out data:

- Learning rate: $\alpha = 0.126$
- Incorrect penalty: $\beta = 0.087$
- Decay rate: $\lambda \approx 0$ (effectively zero)

Algorithm 1 MetaSelector: UCB over Expert Selectors

Require: Base selectors $\mathcal{E} = \{e_1, \dots, e_M\}$, exploration constant c_{meta} , decay rate λ

```
1: Initialize counts  $N_j \leftarrow 0$ , rewards  $R_j \leftarrow 0$  for each expert  $j$ 
2: for  $t = 1, 2, \dots, T$  do
3:   if any expert has  $N_j = 0$  then
4:      $j^* \leftarrow$  any unexplored expert {Round-robin init}
5:   else
6:      $j^* \leftarrow \arg \max_j \frac{R_j}{N_j} + c_{\text{meta}} \sqrt{\frac{\ln \sum_m N_m}{N_j}}$  {UCB over experts}
7:   end if
8:    $i^* \leftarrow e_{j^*}.\text{SELECTARM}(\text{history})$  {Expert picks topic}
9:   Play arm  $i^*$ , observe reward  $r_t$ 
10:  Update student knowledge via Eqs. (1)–(2)
11:   $\bar{k}_t \leftarrow \frac{1}{K} \sum_{i=1}^K k_i^{(t)}$  {Current mean knowledge}
12:   $N_{j^*} \leftarrow N_{j^*} + 1$ 
13:   $R_{j^*} \leftarrow R_{j^*} + \bar{k}_t$  {Credit mean knowledge to expert}
14: end for
```

- Initial knowledge: $k_0 = 0.91$

The near-zero decay rate is consistent with Duolingo’s built-in spaced repetition system already mitigating forgetting. In this regime, most algorithms perform similarly, and the Greedy-Min catastrophe does not manifest strongly.

ASSISTments. We use the ASSISTments 2009–2010 skill-builder dataset [Feng et al., 2009], sampling 500 students with at least 100 interactions, yielding 425,000 total interactions. Knowledge topics correspond to skill tags. Fitted parameters:

- Learning rate: $\alpha = 0.090$
- Incorrect penalty: $\beta = 0.077$
- Decay rate: $\lambda = 0.008$
- Initial knowledge: $k_0 = 0.71$

The non-trivial decay ($\lambda = 0.008$) makes ASSISTments a more informative testbed. BKT-Bandit and PD-Whittle maintain their advantage, and the Greedy-Min catastrophe is visible (Greedy-Min underperforms Random).

S5 Whittle Index Implementation Details

The Whittle advantage index is computed on a discretized state space with $M = 50$ bins over $k \in [0.01, 0.99]$. Value iteration uses convergence threshold 10^{-8} with maximum 500 iterations. The advantage index $W(k) = V_{\text{mixed}}(k) - V_{\text{passive}}(k)$ uses mixed transition probability $1/K$ to account for the budget constraint. Precomputation takes ~ 0.01 seconds per configuration at $M=50$; online lookup is $O(1)$ via nearest-state interpolation.

Sensitivity to discretization. We tested $M \in \{20, 50, 100, 200\}$ and found that coarser grids ($M=20, 50$) produce stable, monotonically decreasing index profiles—high index for low-mastery states, as expected. Finer grids ($M=100, 200$) introduce numerical instabilities in the binary-search/value-iteration pipeline, producing non-monotone index profiles. We therefore use $M=50$ as the default, which balances approximation quality with numerical stability. All granularities compute in under 0.1 seconds.

S6 Extended Limitations

- **Synthetic–real gap.** Our synthetic environment uses a simplified student model with homogeneous learning/forgetting parameters across topics. Real students exhibit topic-dependent learning rates,

prerequisite dependencies, and motivational effects not captured in our simulation. The near-zero decay fitted on Duolingo suggests that real platforms’ existing scheduling may already mitigate the forgetting that drives our synthetic results.

- **Specific failure modes.** BKT-Bandit underperforms at $K=20$, $\lambda=0.005$, likely because its posterior decay mechanism introduces unnecessary uncertainty in near-stationary settings. MetaSelector’s single win ($K=6$, $\lambda=0.01$) indicates that the UCB-over-experts overhead reduces sample efficiency when a single algorithm clearly dominates.
- **Offline evaluation.** Our real-data experiments use offline replay of interaction logs, which introduces distributional shift: the algorithm’s selected arm may differ from the arm that was actually presented, making counterfactual evaluation approximate.
- **Scalability.** The Whittle-advantage algorithms require solving single-arm MDPs via value iteration, which becomes expensive for fine-grained state discretizations. At $K=100$ this adds non-trivial overhead.
- **Fixed hyperparameters.** We use a single exploration constant $c = \sqrt{2}$ for all UCB-family algorithms. Per-algorithm tuning could improve individual performance but would complicate fair comparison.
- **Single-arm activation.** Our RMAB formulation assumes exactly one arm is activated per timestep. Real educational settings may allow multiple questions per session, requiring combinatorial extensions.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: All claims are supported by experiments in Section 4. Performance numbers are reported with standard errors across 3 seeds.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Limitations are discussed in Section 5 and Appendix S6, including the synthetic-real gap and specific failure modes.

3. Theory/Proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (or sketch) proof?

Answer: [NA]

Justification: This is primarily an empirical study. The Greedy-Min catastrophe is demonstrated empirically rather than proved theoretically.

4. Experiments

Question: Does the paper provide sufficient information to reproduce the experiments?

Answer: [Yes]

Justification: Full hyperparameters, environment details, and compute resources are provided in Section 4 and Appendix S1. Code is publicly available.

5. Code and Data

Question: Does the submission provide code and data?

Answer: [Yes]

Justification: Code is included as supplementary material and will be released publicly upon acceptance.

6. Experimental Setting/Details

Question: Does the paper specify all the training and test details?

Answer: [Yes]

Justification: All parameters (α , β , λ , K , number of students, questions) are specified in Section 4 and Appendix S1.

7. Compute

Question: Does the paper report the computational resources used?

Answer: [Yes]

Justification: Compute details (KAUST IBEX cluster, CPU-only, runtimes) are reported in Appendix S1.

8. Error Bars

Question: Does the paper report error bars or confidence intervals?

Answer: [Yes]

Justification: Standard errors across 3 random seeds are reported for all experiments.

9. Broader Impacts

Question: Does the paper discuss potential broader societal impacts?

Answer: [Yes]

Justification: The work targets educational technology. Positive impacts (personalized learning) and risks (over-reliance on automated tutoring) are discussed in Section 5.

10. Safeguards

Question: Does the paper describe safeguards for responsible use?

Answer: [NA]

Justification: The system selects quiz questions and does not pose safety risks.

11. Licenses

Question: Are the creators of assets used in the paper properly credited and license terms respected?

Answer: [Yes]

Justification: Duolingo and ASSISTments datasets are properly cited and used under their respective licenses.

12. New Assets

Question: Are new assets introduced in the paper well documented?

Answer: [Yes]

Justification: Our benchmark code and synthetic environment are documented and will be released with an open-source license.

13. Human Subjects

Question: Does the paper involve crowdsourcing or human subjects?

Answer: [No]

Justification: We use existing public datasets and synthetic simulations only.

14. IRB Approval

Question: Does the research involve human subjects requiring IRB approval?

Answer: [No]

Justification: No new human data was collected.